

PROGRESS UPDATE

Weekly Progress Report

Sentry Agentic - SentinelOne Intelligence Layer

Report Date: July 31, 2026 | Reporting Period: July 27-31, 2026

Overview

This week's work centred on two connected pieces: understanding what actually exists in our SentinelOne environment, and building the data retrieval layer - the tested pathway that lets our AI assistant answer questions using live data instead of guessing or reasoning it out fresh each time.

This mattered because, once environment understanding was established, we ran a scoped test of the assistant that surfaced a real case: a security question answered with false confidence, pulled from the wrong internal source instead of the actual security data - in a security operations context, a mistake like that isn't a cosmetic bug, it's the kind of thing that could mean a missed threat.

To fix this at the root, we built a verified, tested pathway for every type of security question our analysts ask, so the assistant now only answers from a validated, live source, and is designed to say "I don't know" rather than guess when it isn't certain.

Output: a validated retrieval layer that prevents false certainty. This foundational work is now complete and running live - all of the question types our team regularly relies on are covered and passing live accuracy testing, up from effectively zero reliable coverage when we started.

Accomplishments for the Week

Two pieces of work were completed and delivered this week:

Task / Solution	Objectives	Progress Status	Pending Deliverables	Identified Blockers	Timeline
Environment Understanding	Build a verified, live-tested map of what actually exists in the SentinelOne environment - devices, alerts, vulnerabilities, and more - as the foundation everything else is built on.	Completed	None.	None.	Delivered
Data Retrieval Layer (Tested Pathway)	Build a proven, tested way for the assistant to answer the 13 most common analyst questions from live data, replacing guesswork with validated, repeatable retrieval - directly in response to a scoped test finding a false-confidence, wrong-source answer.	Completed	None.	SentinelOne's own AI query-assistant remains unavailable on SentinelOne's side, which limits one question type (deep threat hunting) to hand-built searches. See Notes.	Delivered

Notes - Challenges & Blockers Encountered

Scoped test surfaced a false-confidence answer

Testing the assistant against the newly-mapped environment surfaced a real case: a security question answered confidently from the wrong source. This is the finding that shaped this week's retrieval-layer work (see Overview) and is now resolved.

SentinelOne AI feature unavailable (ongoing)

A SentinelOne-side AI feature used for advanced threat hunting is currently down on SentinelOne's end - not something we can fix directly. It does not block this week's work. We will re-check it before it's needed for Phase 3.

A few smaller technical issues were also found and fixed this week, with no impact to the timeline.

Next Steps for the Week

With this week's foundation complete, next week covers two phases: hardening what's been built under real, live use (Phase 2b), and beginning Agent Capabilities (Phase 3). Planned tasks:

#	Phase	Task	Focus
1	Phase 2b	Fix Live-Use Edge Cases	Harden reliability by closing the edge cases only real, live use surfaces.
2	Phase 2b	Connect to Live Chat	Wire the accuracy layer directly into the live chat experience.
3	Phase 2b	Update Documentation	Bring reference documents in line with this week's changes.
4	Phase 2b	Formal Review & Sign-Off	Review and save this week's work as a completed set of changes.
5	Phase 3	Define Capability Structure	Set the blueprint every capability follows: trigger, plan, inputs, output, and a confidence score.
6	Phase 3	Build the Capability Runner	The engine that runs an investigation step by step, using only proven, validated methods - never improvised.
7	Phase 3	Enforce Safe Composition	Ensure a capability can never invent its own way to fetch data - only reuse what's already proven.
8	Phase 3	Set Up Registry & Test Cases	Create the tracking list of capabilities and a first test case for each, to confirm they work correctly.

Rows 5-8 together make up Phase 3's Framework & Registry stage - the foundation the six specialist roles (Triage, Investigator, Hunter, Correlator, Threat Intelligence, Reporter) are built on. Those roles, plus Session Export and the Daily Briefing, follow in subsequent weeks.

Closing Notes - Insights

Next week's work runs on two fronts: **Post-Hardening**, which turns this week's proof into production reliability, and **Agent Capabilities**, which begins turning that reliability into investigative power analysts can lean on.

Post-Hardening

This is the difference between “works in testing” and “works every time, for every analyst, in the tool they already use.” Concretely: the accuracy layer gets wired directly into live chat, the edge cases only real use surfaces get closed, and this week's work gets formally reviewed and locked in as production-ready. From that point on, every SentinelOne question asked in chat is answered from the same validated, tested source - no exceptions, no regressions.

Agent Capabilities

This is where the assistant evolves from a single, general-purpose responder into a team of specialized sub-agents, each built to do one part of an analyst's job well. Next week we lay the framework and registry that every specialist - triage, investigation, hunting, correlation, reporting - will be built on, composing only from what's already proven, never improvised. It is the deliberate, safe on-ramp to real analytical work, not just faster retrieval.

Together, these two workstreams move the project from trustworthy to operationally valuable - the natural next step in a build sequence that has, so far, proven every claim before shipping it.