

PROJECT TIMELINE

Complete Project History, From Setup to Today

Sentry Agentic - SentinelOne Intelligence Layer

Management Report - 2026-07-31

Overview

This report lists every phase of work on the SentinelOne project, from the platform's original setup through to what's happening next. Each phase only began once the phase before it was tested and confirmed working - nothing here was built in parallel or skipped ahead.

- Dates shown are real work dates, not estimates.
- Phases that haven't been scheduled yet are marked **Not yet scheduled** rather than given a made-up date.
- The list below is complete and in order - nothing is skipped between the platform's original setup and today.

When	Phase	What Was Delivered	Status
Before mid-July 2026	Getting the Platform Running	The core system was built and switched on: the database, the caching layer, the backend service, and the web interface analysts use. Connections to 39 different security tools were configured. The broader team of 13 AI specialist roles used across the whole platform was established.	COMPLETE
Mid-July 2026	Stabilizing the Platform	Three real problems were found and fixed while stabilizing the platform after a rebrand: a mismatch in database login settings that could have pointed the system at the wrong database, a bug where one failed connection could crash the entire system on startup, and a bug where AI provider access keys could be lost every time the system restarted.	COMPLETE - one item needs a follow-up check
Late July 2026	Connecting to SentinelOne	The connection to SentinelOne was established and fixed: a broken internal reference was corrected, the correct version of SentinelOne's official connector was locked in, the access key was placed in secure storage, and the first version of the setup documentation was written.	COMPLETE - later updated, see below

When	Phase	What Was Delivered	Status
Late July 2026	Mapping What SentinelOne Can Do	Every one of SentinelOne's 33 available data tools was catalogued. The project's working structure was set up. A missing planning document that was blocking the next phase was tracked down and added.	COMPLETE
Late July 2026	Exploring the Live Environment	A safe, read-only scan of the real SentinelOne environment was run. It found 385 live alerts and mapped out the full account, site, and group structure. One incorrect assumption, found because it didn't match what was actually seen, was corrected.	COMPLETE
Late July 2026	Defining What Exists	10 real categories of information in this SentinelOne environment - devices, alerts, and so on - were confirmed using real samples, not assumptions. Six categories that don't apply to this environment were correctly ruled out, with the reasoning documented.	COMPLETE
Late July 2026	Building the Question Checklist	A checklist of 12 common question types was built, based only on confirmed facts from the previous steps. Five of the most important were flagged as priority. Every entry was marked with the one data source it must never be answered from.	COMPLETE
Late July 2026	Proving Each Answer Works	A tested, working method for answering 7 of the question types was built and confirmed against the real environment, including what happens when there's simply no data to return.	COMPLETE
Late July 2026	Deep Investigation Tools	A reference guide and 6 ready-made investigation patterns, tied to known attack techniques, were built by hand after confirming that SentinelOne's own AI-assisted query builder is currently unavailable. One bug in the matching logic was found and fixed.	COMPLETE - one related SentinelOne feature is unavailable
Late July 2026	Teaching the System to Recognize Questions	A private, self-contained matching system was built so the system can tell which question type a new question belongs to, without sending anything outside the system to do it. It was tuned until it recognized real-world phrasing reliably.	COMPLETE

When	Phase	What Was Delivered	Status
Late July 2026	Building the Accuracy Check	The safeguard that checks every answer before it's given was built: it checks the source is right, checks whether a result is genuinely empty versus simply not found, and refuses to guess. Testing against a larger, more realistic sample of alerts found and fixed a serious undercount in a key vulnerability figure - the system had been reporting 1 when the real number was nearly 38,000.	COMPLETE
Late July 2026	Final Testing	Every one of the 12 question types was tested against the real, live environment and passed. Two small but real bugs were found and fixed during this final round, including a mislabeled data source in one of the answer formats.	COMPLETE
Late July 2026	Fixing Issues Found in Live Use	Two real problems, only visible once real people started using the system, were found and fixed: a question about "the last 24 hours" was being silently treated as "all time," and a natural follow-up question about total device count had no way to be answered correctly. A 13th question type was added to fix the second issue. The whole system was then connected directly into the live chat experience for the first time, and the original setup documentation was corrected to match everything learned.	COMPLETE - two documentation updates still to be made
Not yet scheduled	Building the Investigation Framework	Define exactly how a full investigation is structured, build the engine that runs one step by step, and make sure it can only ever use proven, tested methods - never improvise.	NOT STARTED
Not yet scheduled	Triage	Teach the system to judge whether a threat is real, how serious it is, and how far it may have spread.	NOT STARTED
Not yet scheduled	Investigator	Teach the system to reconstruct a full attack chain across devices and map it to known attack techniques.	NOT STARTED
Not yet scheduled	Hunter	Run scheduled, proactive searches for known attack patterns and report anything found.	NOT STARTED
Not yet scheduled	Correlator	Spot patterns that connect separate alerts into one bigger picture, clearly separating fact from interpretation.	NOT STARTED
Not yet scheduled	Threat Intelligence	Automatically enrich a suspicious file or indicator with reputation and context from trusted sources.	NOT STARTED
Not yet scheduled	Reporter and Exports	Assemble full, audit-ready reports, and let any investigation be downloaded as a PDF, Word document, or similar.	NOT STARTED

When	Phase	What Was Delivered	Status
Not yet scheduled	Daily Briefing	Produce an automatic daily summary of what mattered in the environment over the last 24 hours - reporting only, never taking action.	NOT STARTED
Not yet scheduled	Final Testing (Investigations)	Test every one of the above against real data before considering this phase complete.	NOT STARTED
Not yet scheduled	Letting the System Take Action	Decide which support-ticketing system to connect to - not yet decided. Build the ability to recommend and eventually create tickets. Connect it to the existing approval system so a person signs off before anything happens. Run it under close supervision before allowing anything fully automatic.	NOT STARTED

Conclusion

This history shows a carefully sequenced build, not trial and error: each phase's checks passed before the next one began, and real problems were found and fixed at nearly every step rather than glossed over - 3 problems in the underlying platform before the SentinelOne work even began, and 12 more within the SentinelOne work itself. Every one of them is explained, with its cause and its fix, in the Progress Update.

As of 2026-07-31: the goal of getting accurate, checkable answers from SentinelOne is complete and is being used in the live product - all 13 question types the system was tested against now answer correctly. Teaching the system to run full investigations, and eventually letting it take action on its own, are both fully planned but haven't started. Both are ready to begin as soon as they're scheduled, and - for letting the system take action - once a decision is made on which ticketing system to use. See the Progress Update for the full list of what's still open and what happens next.